

1. Which of the following best describes the main goal of tokenization in NLP?
 - A. To label each token with a grammatical category
 - B. To convert text into a sequence of symbols suitable for modeling
 - C. To reduce vocabulary size to zero
 - D. To assign probabilities to sentences
2. Why is splitting text by spaces (`text.split(" ")`) considered an insufficient tokenizer?
 - A. It produces too many tokens
 - B. It does not work well for languages without explicit word boundaries
 - C. It cannot tokenize English at all
 - D. It ignores punctuation entirely
3. Which of the following is NOT listed as a common sequence tagging task?
 - A. Part-of-speech tagging
 - B. Named Entity Recognition
 - C. Dependency parsing
 - D. Sentiment analysis
4. What is a common limitation of BIO tagging for Named Entity Recognition?
 - A. It cannot represent overlapping or nested entities
 - B. It requires too many labels
 - C. It does not support multiple entity types
 - D. It ignores entity boundaries
5. When choosing between character-level tokenization and sentence-level tokenization, which trade-off must be considered?
 - A. Speed vs accuracy
 - B. Tokenized sequence length vs vocabulary size
 - C. Precision vs recall
 - D. Memory vs computation

Answers:

1) B 2)B 3)D 4)A 5)B