



CSE-476: Introduction to Natural Language Processing
Quiz 7

Multiple Choice Questions (Single correct)

Instructions: Select the correct option with a short explanation.

- Q1. What is a key limitation of a simple Retrieval-Augmented Generation (RAG) system?
- (A) It cannot process text input
 - (B) It struggles with noisy or irrelevant retrieved data
 - (C) It cannot use neural networks
 - (D) It only works for images
- Q2. Which situation would most likely produce a high MRR but relatively low MAP?
- (A) Relevant documents are evenly distributed across rankings
 - (B) The first relevant document appears early, but others are ranked poorly
 - (C) All relevant documents appear at the top
 - (D) No relevant documents are retrieved
- Q3. Why does Mean Average Precision provide a more comprehensive evaluation than Precision@k?
- (A) It accounts for ranking positions of all relevant documents
 - (B) It only considers the first relevant document
 - (C) It ignores irrelevant documents
 - (D) It requires fewer computations
- Q4. In Dense Passage Retrieval, what is the primary role of negative samples during training?
- (A) To enforce similarity between unrelated queries and documents
 - (B) To push irrelevant documents away in embedding space
 - (C) To reduce model complexity
 - (D) To increase dataset size
- Q5. Which statement best explains why BM25 may outperform DPR on unseen domains?
- (A) DPR requires labeled training data and may not generalize well

- (B) BM25 uses neural embeddings
 - (C) BM25 ignores word frequencies
 - (D) DPR does not use queries
- Q6. What is the key reason self-consistency requires a non-zero temperature during inference?
- (A) To produce diverse reasoning paths
 - (B) To reduce computation time
 - (C) To eliminate randomness
 - (D) To ensure deterministic outputs
- Q7. During DPR training, what happens if negative samples are too easy (clearly unrelated)?
- (A) The model may fail to learn fine-grained differences
 - (B) The model learns stronger distinctions
 - (C) The training becomes unstable
 - (D) The embeddings collapse
- Q8. Which issue arises in RAG when retrieved documents contain conflicting information?
- (A) Improved model confidence
 - (B) Reduced inference time
 - (C) Automatic conflict resolution
 - (D) Increased hallucination risk
- Q9. Why does increasing Chain-of-Thought (CoT) length sometimes improve reasoning?
- (A) It reduces randomness
 - (B) It allows more intermediate reasoning steps
 - (C) It eliminates errors
 - (D) It speeds up inference
- Q10. What is the key idea behind test-time scaling?
- (A) Train larger models
 - (B) Increase dataset size
 - (C) Extend reasoning at inference time
 - (D) Reduce token usage

End of Paper